

1. Input Libraries

```
import pandas as pd
import numpy as np

from sklearn.linear_model import LinearRegression , Ridge, Lasso
from sklearn.model_selection import train_test_split

from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

from sklearn.preprocessing import OneHotEncoder,LabelEncoder

from statsmodels.stats.outliers_influence import variance_inflation_factor

import matplotlib.pyplot as plt
import seaborn as sns

from scipy.stats import shapiro , kstest, normaltest, skew
```

2. Problem statement

To predict FISH WEIGHT by using below features features = ['Species', 'Length1', 'Length2', 'Length3', 'Height', 'Width'] Target = 'Weight'

3. Data Gathering

In [144]:

```
raw_df = pd.read_csv("Fish.csv")
raw_df.head()
```

Out[144]:

	Species	Weight	Length1	Length2	Length3	Height	Width
0	Bream	242.0	23.2	25.4	30.0	11.5200	4.0200
1	Bream	290.0	24.0	26.3	31.2	12.4800	4.3056
2	Bream	340.0	23.9	26.5	31.1	12.3778	4.6961
3	Bream	363.0	26.3	29.0	33.5	12.7300	4.4555
4	Bream	430.0	26.5	29.0	34.0	12.4440	5.1340

4,5,6 EDA (Exploratory Data Analysis), Feature Engineering & Feature selection

In [52]:

```
raw_df.isna().sum() # checked if there are Null or missing values
```

Out[52]:

```
Species      0
Weight       0
Length1      0
Length2      0
Length3      0
Height       0
Width        0
dtype: int64
```

In [53]:

```
raw_df.info() # checked if any null value is represented by any symbol and what is data type
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 159 entries, 0 to 158
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   Species    159 non-null    object
 1   Weight     159 non-null    float64
 2   Length1    159 non-null    float64
 3   Length2    159 non-null    float64
 4   Length3    159 non-null    float64
 5   Height     159 non-null    float64
 6   Width      159 non-null    float64
dtypes: float64(6), object(1)
memory usage: 8.8+ KB
```

1. Linearity

In [54]:

```
raw_df.corr()
```

Out[54]:

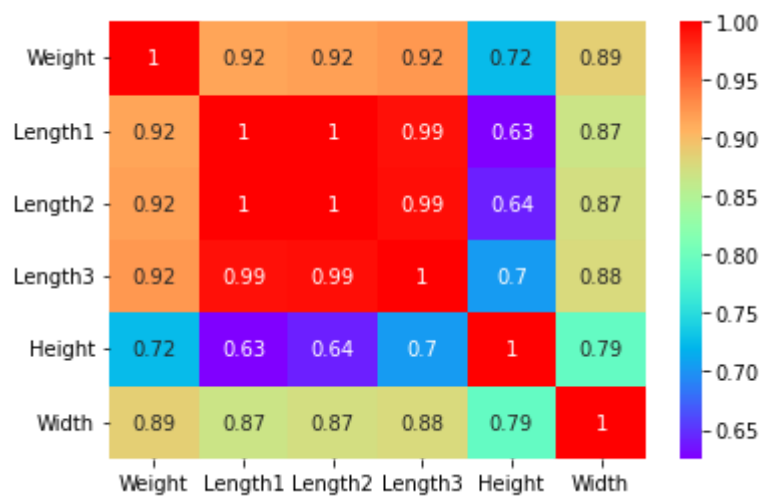
	Weight	Length1	Length2	Length3	Height	Width
Weight	1.000000	0.915712	0.918618	0.923044	0.724345	0.886507
Length1	0.915712	1.000000	0.999517	0.992031	0.625378	0.867050
Length2	0.918618	0.999517	1.000000	0.994103	0.640441	0.873547
Length3	0.923044	0.992031	0.994103	1.000000	0.703409	0.878520
Height	0.724345	0.625378	0.640441	0.703409	1.000000	0.792881
Width	0.886507	0.867050	0.873547	0.878520	0.792881	1.000000

In [55]:

```
sns.heatmap(raw_df.corr(),annot = True, cmap = 'rainbow')
```

Out[55]:

<AxesSubplot:>



2. MultiColinearity

In [56]:

```

vif_list = []
for i in range(raw_df.shape[1]):
    vif = variance_inflation_factor(raw_df.to_numpy(), i)
    vif_list.append(vif)

s1 = pd.Series(vif_list, index = raw_df.columns)    # converted into series
s1.sort_values().plot(kind = 'barh')

```

TypeError

Traceback (most recent call last)

Input In [56], in <cell line: 2>()

```

1 vif_list = []
2 for i in range(raw_df.shape[1]):
----> 3     vif = variance_inflation_factor(raw_df.to_numpy(), i)
4     vif_list.append(vif)
6 s1 = pd.Series(vif_list, index = raw_df.columns)    # converted into series
o series

```

File ~\AppData\Local\Programs\Python\Python310\lib\site-packages\statsmodels\stats\outliers_influence.py:194, in variance_inflation_factor(exog, exog_idx x)

```

192 mask = np.arange(k_vars) != exog_idx
193 x_noti = exog[:, mask]
--> 194 r_squared_i = OLS(x_i, x_noti).fit().rsquared
195 vif = 1. / (1. - r_squared_i)
196 return vif

```

File ~\AppData\Local\Programs\Python\Python310\lib\site-packages\statsmodels\regression\linear_model.py:890, in OLS.__init__(self, endog, exog, missing, hasconst, **kwargs)

```

887 msg = ("Weights are not supported in OLS and will be ignored"
888        "An exception will be raised in the next version.")
889 warnings.warn(msg, ValueWarning)
--> 890 super(OLS, self).__init__(endog, exog, missing=missing,
891                             hasconst=hasconst, **kwargs)
892 if "weights" in self._init_keys:
893     self._init_keys.remove("weights")

```

File ~\AppData\Local\Programs\Python\Python310\lib\site-packages\statsmodels\regression\linear_model.py:717, in WLS.__init__(self, endog, exog, weights, missing, hasconst, **kwargs)

```

715 else:
716     weights = weights.squeeze()
--> 717 super(WLS, self).__init__(endog, exog, missing=missing,
718                             weights=weights, hasconst=hasconst, **kwargs)
719 nobs = self.exog.shape[0]
720 weights = self.weights

```

File ~\AppData\Local\Programs\Python\Python310\lib\site-packages\statsmodels\regression\linear_model.py:191, in RegressionModel.__init__(self, endog, exog, **kwargs)

```

190 def __init__(self, endog, exog, **kwargs):
--> 191     super(RegressionModel, self).__init__(endog, exog, **kwargs)
192     self._data_attr.extend(['pinv_wexog', 'wendog', 'wexog', 'weights'])

```

File ~\AppData\Local\Programs\Python\Python310\lib\site-packages\statsmodels\base\model.py:267, in LikelihoodModel.__init__(self, endog, exog, **kwargs)

```

266 def __init__(self, endog, exog=None, **kwargs):
--> 267     super().__init__(endog, exog, **kwargs)
268     self.initialize()

```

File ~\AppData\Local\Programs\Python\Python310\lib\site-packages\statsmodels\base\model.py:92, in Model.__init__(self, endog, exog, **kwargs)

```

90 missing = kwargs.pop('missing', 'none')
91 hasconst = kwargs.pop('hasconst', None)
--> 92 self.data = self._handle_data(endog, exog, missing, hasconst,
93                                **kwargs)
94 self.k_constant = self.data.k_constant
95 self.exog = self.data.exog

```

File ~\AppData\Local\Programs\Python\Python310\lib\site-packages\statsmodels\base\model.py:132, in Model._handle_data(self, endog, exog, missing, hasconst, **kwargs)

```

131 def _handle_data(self, endog, exog, missing, hasconst, **kwargs):
--> 132     data = handle_data(endog, exog, missing, hasconst, **kwargs)
133     # kwargs arrays could have changed, easier to just attach here
134     for key in kwargs:

```

File ~\AppData\Local\Programs\Python\Python310\lib\site-packages\statsmodels\base\data.py:673, in handle_data(endog, exog, missing, hasconst, **kwargs)

```

670     exog = np.asarray(exog)
672 klass = handle_data_class_factory(endog, exog)
--> 673 return klass(endog, exog=exog, missing=missing, hasconst=hasconst,
674               **kwargs)

```

File ~\AppData\Local\Programs\Python\Python310\lib\site-packages\statsmodels\base\data.py:86, in ModelData.__init__(self, endog, exog, missing, hasconst, **kwargs)

```

84 self.const_idx = None
85 self.k_constant = 0
--> 86 self._handle_constant(hasconst)
87 self._check_integrity()
88 self._cache = {}

```

File ~\AppData\Local\Programs\Python\Python310\lib\site-packages\statsmodels\base\data.py:131, in ModelData._handle_constant(self, hasconst)

```

129 check_implicit = False
130 exog_max = np.max(self.exog, axis=0)
--> 131 if not np.isfinite(exog_max).all():
132     raise MissingDataError('exog contains inf or nans')
133 exog_min = np.min(self.exog, axis=0)

```

TypeError: ufunc 'isfinite' not supported for the input types, and the inputs could not be safely coerced to any supported types according to the casting rule 'safe'

In [57]:

```
vif = [variance_inflation_factor(raw_df.tonumpy(), i) for i in range(raw_df.shape[1])]
```

AttributeError

Traceback (most recent call last)

Input In [57], in <cell line: 1>()

```
----> 1 vif = [variance_inflation_factor(raw_df.tonumpy(), i) for i in range
(raw_df.shape[1])]
```

Input In [57], in <listcomp>(.0)

```
----> 1 vif = [variance_inflation_factor(raw_df.tonumpy(), i) for i in range
(raw_df.shape[1])]
```

File ~\AppData\Local\Programs\Python\Python310\lib\site-packages\pandas\core\generic.py:5575, in NDFrame.__getattr__(self, name)

```
5568 if (
5569     name not in self._internal_names_set
5570     and name not in self._metadata
5571     and name not in self._accessors
5572     and self._info_axis._can_hold_identifiers_and_holds_name(name)
5573 ):
5574     return self[name]
-> 5575 return object.__getattr__(self, name)
```

AttributeError: 'DataFrame' object has no attribute 'tonumpy'

In [58]:

```
raw_df.to_numpy()
```

Out[58]:

```
array([[ 'Bream', 242.0, 23.2, ..., 30.0, 11.52, 4.02],
       [ 'Bream', 290.0, 24.0, ..., 31.2, 12.48, 4.3056],
       [ 'Bream', 340.0, 23.9, ..., 31.1, 12.3778, 4.6961],
       ...,
       [ 'Smelt', 12.2, 12.1, ..., 13.8, 2.277, 1.2558],
       [ 'Smelt', 19.7, 13.2, ..., 15.2, 2.8728, 2.0672],
       [ 'Smelt', 19.9, 13.8, ..., 16.2, 2.9322, 1.8792]], dtype=object)
```

Detecting Outliers in each column

In [59]:

```
raw_df['Weight']
```

Out[59]:

```
0      242.0
1      290.0
2      340.0
3      363.0
4      430.0
...
154     12.2
155     13.4
156     12.2
157     19.7
158     19.9
```


Name: Weight, Length: 159, dtype: float64

In [60]:

```
sns.kdeplot(raw_df['Weight'])
```

Out[60]:

<AxesSubplot:xlabel='Weight', ylabel='Density'>



In [61]:

```
q1 = raw_df['Weight'].quantile(0.25)
q2 = raw_df['Weight'].quantile(0.50)
q3 = raw_df['Weight'].quantile(0.75)

IQR = q3 - q1

upper_tail = q3 + 1.5*IQR
lower_tail = q1 - 1.5*IQR

median = raw_df['Weight'].median()

print('q1', q1)
print('q2', q2)
print('q3', q3)
print("Median", median)
print('upper_tail', upper_tail)
print('lower_tail', lower_tail)
```

```
q1 120.0
q2 273.0
q3 650.0
Median 273.0
upper_tail 1445.0
lower_tail -675.0
```

In [62]:


```
sns.boxplot(raw_df['Weight'])
```

C:\Users\ADMIN\AppData\Local\Programs\Python\Python310\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[62]:

<AxesSubplot:xlabel='Weight'>



In [63]:

```
raw_df[['Weight']].loc[raw_df['Weight'] > upper_tail ]
```

Out[63]:

	Weight
142	1600.0
143	1550.0
144	1650.0

In [64]:

```
raw_df.loc[raw_df['Weight'] > upper_tail, 'Weight'] = upper_tail  
print(raw_df[['Weight']].loc[raw_df['Weight'] > upper_tail ])
```

Empty DataFrame
Columns: [Weight]
Index: []

In [65]:


```
sns.boxplot(raw_df['Weight'])
```

C:\Users\ADMIN\AppData\Local\Programs\Python\Python310\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[65]:

<AxesSubplot:xlabel='Weight'>



In [72]:

```
raw_df['Length1']
```

Out[72]:

0 23.2

1 24.0

2 23.9

3 26.3

4 26.5

...

154 11.5

155 11.7

156 12.1

157 13.2

158 13.8


Name: Length1, Length: 159, dtype: float64

In [73]:

```
sns.kdeplot(raw_df['Length1'])
```

Out[73]:

<AxesSubplot:xlabel='Length1', ylabel='Density'>



In [74]:


```
sns.boxplot(raw_df['Length1'])
```

C:\Users\ADMIN\AppData\Local\Programs\Python\Python310\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[74]:

<AxesSubplot:xlabel='Length1'>



In [75]:

```

q1 = raw_df['Length1'].quantile(0.25)
q2 = raw_df['Length1'].quantile(0.50)
q3 = raw_df['Length1'].quantile(0.75)

IQR = q3 - q1

upper_tail = q3 + 1.5*IQR
lower_tail = q1 - 1.5*IQR

median = raw_df['Length1'].median()

print('q1', q1)
print('q2', q2)
print('q3', q3)
print("Median", median)
print('upper_tail', upper_tail)
print('lower_tail', lower_tail)

```

```

q1 19.05
q2 25.2
q3 32.7
Median 25.2
upper_tail 53.175000000000004
lower_tail -1.4250000000000007

```

In [78]:

```
raw_df[['Length1']].loc[raw_df['Length1'] > upper_tail]
```

Out[78]:

	Length1
142	56.0
143	56.0
144	59.0

In [79]:

```

raw_df.loc[raw_df['Length1'] > upper_tail, 'Length1'] = upper_tail
print(raw_df[['Length1']].loc[raw_df['Length1'] > upper_tail ])

```

```

Empty DataFrame
Columns: [Length1]
Index: []

```

In [80]:

```
raw_df['Length2']
```

Out[80]:

0 25.4

1 26.3

2 26.5

3 29.0

4 29.0

...

154 12.2

155 12.4

156 13.0

157 14.3

158 15.0


Name: Length2, Length: 159, dtype: float64

In [81]:

```
sns.kdeplot(raw_df['Length2'])
```

Out[81]:

<AxesSubplot:xlabel='Length2', ylabel='Density'>



In [82]:


```
sns.boxplot(raw_df['Length2'])
```

C:\Users\ADMIN\AppData\Local\Programs\Python\Python310\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[82]:

<AxesSubplot:xlabel='Length2'>



In [83]:

```
q1 = raw_df['Length2'].quantile(0.25)
q2 = raw_df['Length2'].quantile(0.50)
q3 = raw_df['Length2'].quantile(0.75)
```

```
IQR = q3 - q1
```

```
upper_tail = q3 + 1.5*IQR
```

```
lower_tail = q1 - 1.5*IQR
```

```
median = raw_df['Length2'].median()
```

```
print('q1', q1)
print('q2', q2)
print('q3', q3)
print("Median", median)
print('upper_tail', upper_tail)
print('lower_tail', lower_tail)
```

```
q1 21.0
```

```
q2 27.3
```

```
q3 35.5
```

```
Median 27.3
```

```
upper_tail 57.25
```

```
lower_tail -0.75
```

In [84]:

```
raw_df[['Length2']].loc[raw_df['Length2'] > upper_tail]
```

Out[84]:

	Length2
142	60.0
143	60.0
144	63.4

In [85]:

```
raw_df.loc[raw_df['Length2'] > upper_tail, 'Length2'] = upper_tail
print(raw_df[['Length2']].loc[raw_df['Length2'] > upper_tail ])
```

Empty DataFrame
Columns: [Length2]
Index: []

In [86]:

```
raw_df['Length3']
```

Out[86]:

0	30.0
1	31.2
2	31.1
3	33.5
4	34.0
	...
154	13.4
155	13.5
156	13.8
157	15.2
158	16.2


Name: Length3, Length: 159, dtype: float64

In [89]:

```
sns.kdeplot(raw_df['Length3'])
```

Out[89]:

<AxesSubplot:xlabel='Length3', ylabel='Density'>



In [90]:


```
sns.boxplot(raw_df['Length3'])
```

C:\Users\ADMIN\AppData\Local\Programs\Python\Python310\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[90]:

<AxesSubplot:xlabel='Length3'>



In [91]:

```

q1 = raw_df['Length3'].quantile(0.25)
q2 = raw_df['Length3'].quantile(0.50)
q3 = raw_df['Length3'].quantile(0.75)

IQR = q3 - q1

upper_tail = q3 + 1.5*IQR
lower_tail = q1 - 1.5*IQR

median = raw_df['Length3'].median()

print('q1', q1)
print('q2', q2)
print('q3', q3)
print("Median", median)
print('upper_tail', upper_tail)
print('lower_tail', lower_tail)

```

```

q1 23.15
q2 29.4
q3 39.650000000000006
Median 29.4
upper_tail 64.40000000000002
lower_tail -1.600000000000012

```

In [93]:

```
raw_df[['Length3']].loc[raw_df['Length3'] > upper_tail]
```

Out[93]:

	Length3
144	68.0

In [94]:

```

raw_df.loc[raw_df['Length2'] > upper_tail, 'Length2'] = upper_tail
print(raw_df[['Length2']].loc[raw_df['Length2'] > upper_tail ])

```

```

Empty DataFrame
Columns: [Length2]
Index: []

```

In [95]:

```
raw_df['Width']
```

Out[95]:

```
0      4.0200  
1      4.3056  
2      4.6961  
3      4.4555  
4      5.1340
```

...

```
154    1.3936  
155    1.2690  
156    1.2558  
157    2.0672  
158    1.8792
```


Name: Width, Length: 159, dtype: float64

In [96]:

```
sns.kdeplot(raw_df['Width'])
```

Out[96]:

<AxesSubplot:xlabel='Width', ylabel='Density'>



In [97]:


```
sns.boxplot(raw_df['Width'])
```

C:\Users\ADMIN\AppData\Local\Programs\Python\Python310\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[97]:

<AxesSubplot:xlabel='Width'>



In [98]:

```
raw_df['Height']
```

Out[98]:

```
0      11.5200
1      12.4800
2      12.3778
3      12.7300
4      12.4440
```

```
...
```

```
154     2.0904
155     2.4300
156     2.2770
157     2.8728
158     2.9322
```


Name: Height, Length: 159, dtype: float64

In [99]:

```
sns.kdeplot(raw_df['Height'])
```

Out[99]:

<AxesSubplot:xlabel='Height', ylabel='Density'>



In [100]:


```
sns.boxplot(raw_df['Height'])
```

C:\Users\ADMIN\AppData\Local\Programs\Python\Python310\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[100]:

<AxesSubplot:xlabel='Height'>



In [101]:

```
raw_df['Species']
```

Out[101]:

```
0    Bream
1    Bream
2    Bream
3    Bream
4    Bream
```

```
...
154  Smelt
155  Smelt
156  Smelt
157  Smelt
158  Smelt
```

Name: Species, Length: 159, dtype: object

In [121]:

```
raw_df['Species'].unique().tolist()
```

Out[121]:

```
['Bream', 'Roach', 'Whitefish', 'Parkki', 'Perch', 'Pike', 'Smelt']
```

In [145]:

```
raw_df = pd.get_dummies(raw_df, columns=['Species'])
raw_df
```

Out[145]:

	Weight	Length1	Length2	Length3	Height	Width	Species_Bream	Species_Parkki	Spe
0	242.0	23.2	25.4	30.0	11.5200	4.0200	1	0	
1	290.0	24.0	26.3	31.2	12.4800	4.3056	1	0	
2	340.0	23.9	26.5	31.1	12.3778	4.6961	1	0	
3	363.0	26.3	29.0	33.5	12.7300	4.4555	1	0	
4	430.0	26.5	29.0	34.0	12.4440	5.1340	1	0	
...
154	12.2	11.5	12.2	13.4	2.0904	1.3936	0	0	
155	13.4	11.7	12.4	13.5	2.4300	1.2690	0	0	
156	12.2	12.1	13.0	13.8	2.2770	1.2558	0	0	
157	19.7	13.2	14.3	15.2	2.8728	2.0672	0	0	
158	19.9	13.8	15.0	16.2	2.9322	1.8792	0	0	

159 rows × 13 columns

Train test split

In [146]:

```
x = raw_df.drop('Weight', axis = 1)
y = raw_df['Weight']
raw_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 159 entries, 0 to 158
Data columns (total 13 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Weight              159 non-null    float64
1   Length1             159 non-null    float64
2   Length2             159 non-null    float64
3   Length3             159 non-null    float64
4   Height              159 non-null    float64
5   Width               159 non-null    float64
6   Species_Bream       159 non-null    uint8
7   Species_Parkki      159 non-null    uint8
8   Species_Perch       159 non-null    uint8
9   Species_Pike        159 non-null    uint8
10  Species_Roach       159 non-null    uint8
11  Species_Smelt       159 non-null    uint8
12  Species_Whitefish   159 non-null    uint8
dtypes: float64(6), uint8(7)
memory usage: 8.7 KB
```

In [147]:

```
x_train, x_test, y_train, y_test = train_test_split(x,y, test_size = 0.2, random_state = 20)
```

In [148]:

```
# Instantiating Linear Regression
```

```
linear_reg_model = LinearRegression()
linear_reg_model.fit(x_train, y_train) # Gradient Descent >> best m and c values
```

Out[148]:

```
▼ LinearRegression
LinearRegression()
```

Model Evaluation on Testing Data

In [151]:

```
## Model Evaluation on Testing Data
```

```
y_pred = linear_reg_model.predict(x_test)
```

```
mse = mean_squared_error(y_test, y_pred)
print("Mean Squared Error :",mse)
```

```
rmse = np.sqrt(mse)
print("Root Mean Squared Error :",rmse)
```

```
mae = mean_absolute_error(y_test, y_pred)
print("Mean Absolute Error :",mae)
```

```
r_squared_value = r2_score(y_test, y_pred)
print("R Squared Value :",r_squared_value)
```

```
adj_r2 = 1 - (((1-r_squared_value) * (x_test.shape[0] - 1))/(x_test.shape[0] - x_test.shape[1]))
print("Adjusted R-Squared Value :",adj_r2)
```

```
r2 = linear_reg_model.score(x_test, y_test) # Without Predict Function
print("R2 :",r2)
```

```
Mean Squared Error : 10641.346482529983
Root Mean Squared Error : 103.1569022534604
Mean Absolute Error : 76.334073462639
R Squared Value : 0.9241345503088414
Adjusted R-Squared Value : 0.8762195294512676
R2 : 0.9241345503088414
```

Model Evaluation on Training Data

In [150]:

```
## Model Evaluation on Training Data
```

```
y_pred_train = linear_reg_model.predict(x_train) # 404 rows
```

```
mse = mean_squared_error(y_train, y_pred_train)  
print("Mean Squared Error :",mse)
```

```
rmse = np.sqrt(mse)  
print("Root Mean Squared Error :",rmse)
```

```
mae = mean_absolute_error(y_train, y_pred_train)  
print("Mean Absolute Error :",mae)
```

```
r_squared_value = r2_score(y_train, y_pred_train)  
print("R Squared Value :",r_squared_value)
```

```
adj_r2 = 1 - (((1-r_squared_value) * (x_train.shape[0] - 1))/(x_train.shape[0] - x_train.sha  
print("Adjusted R-Squared Value :",adj_r2)
```

```
r2 = linear_reg_model.score(x_train, y_train)  
print("R2 :",r2)
```

Mean Squared Error : 7657.374223249367

Root Mean Squared Error : 87.50642389704522

Mean Absolute Error : 63.392048561968075

R Squared Value : 0.9382732908398195

Adjusted R-Squared Value : 0.9317757425071689

R2 : 0.9382732908398195

Testing on Single Row

In [153]:

```
x.head(1).T
```

Out[153]:

	0
Length1	23.20
Length2	25.40
Length3	30.00
Height	11.52
Width	4.02
Species_Bream	1.00
Species_Parkki	0.00
Species_Perch	0.00
Species_Pike	0.00
Species_Roach	0.00
Species_Smelt	0.00
Species_Whitefish	0.00

In [167]:

```
Length1 = 23.20
Length2 = 25.40
Length3 = 30.00
Height = 11.52
Width = 4.02
Species = 'Parkki'
```

```
# Weight = ?
```

In [168]:

```
project_data = {'columns' : list(x.columns)}
```

In [169]:

```
column_names = x.columns
column_names
```

Out[169]:

```
Index(['Length1', 'Length2', 'Length3', 'Height', 'Width', 'Species_Bream',
      'Species_Parkki', 'Species_Perch', 'Species_Pike', 'Species_Roach',
      'Species_Smelt', 'Species_Whitefish'],
      dtype='object')
```

In [170]:

```
Species = 'Species_' + Species
Species_index = np.where(column_names == Species)
```

In [171]:

```
test_array = np.zeros(x.shape[1])
test_array[0] = Length1
test_array[1] = Length2
test_array[2] = Length3
test_array[3] = Height
test_array[4] = Width
test_array[Species_index] = 1

test_array
```

Out[171]:

```
array([23.2 , 25.4 , 30.   , 11.52,  4.02,  0.   ,  1.   ,  0.   ,  0.   ,
        0.   ,  0.   ,  0.   ])
```

In [173]:

```
linear_reg_model.predict([test_array])
```

```
C:\Users\ADMIN\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names, but LinearRegression was fitted with feature names
  warnings.warn(
```

Out[173]:

```
array([444.90795215])
```